ρlατform

Platform v2 - Press Release

Subnet 100, Bittensor

Cortex Foundation

Subnet 100, Bittensor

February 20, 2026

---

1 Platform v2: A New Chapter for Subnet 100

February 20, 2026 - Cortex Foundation

Over the past several weeks, the Platform Network team has conducted extensive research into both the development of our software products and the evolution of the subnet itself. Today, we are announcing a series of major upgrades that mark a turning point for Subnet 100. This document covers the Platform v2 infrastructure upgrade, the overhaul of our incentive system, the alpha launch of Cortex IDE, and the measures we are taking to ensure that every miner on this network produces real, measurable value.

2 Platform v2 — Automatic P2P Upgrade

We are proceeding with the automatic upgrade of Platform v2 across all validators. This upgrade transitions the entire subnet to a fully decentralized peer-to-peer architecture. Validators will now form a libp2p mesh, execute challenge logic inside a hardened WASM runtime, reach stake-weighted consensus, and submit finalized weights to the Bittensor chain. There is no centralized relay. Every validator independently verifies miner submissions and participates in consensus.

This upgrade is automatic. All validators will receive the new binary and transition seamlessly. No manual intervention is required from validator operators.

The full technical specification of the P2P architecture, consensus protocol, and WASM runtime is available in our roadmap document.

3 Term Challenge v2 — Ending Static Benchmark Gaming

We suspended Term Challenge emissions one week ago. This was a deliberate decision.

We observed that miners were hardcoding their agents to pass the specific tasks in the Terminal Benchmark 2.0 dataset. Rather than building genuinely capable software engineering agents, miners were memorizing solutions and gaming the static evaluation set. This is not viable. It produces no value for the network, no improvement in the software, and no useful training data. We decided to put an end to it.

During this week of suspension, we focused our efforts on refining swe-forge, our synthetic dataset generation tool.

swe-forge mines GitHub Archive for recently merged pull requests, enriches them with repository metadata, classifies difficulty via LLM, and generates test cases inside Docker-isolated containers through a multi-turn agentic loop. Every generated task is validated through a dual-commit process: tests must fail on the pre-PR commit and pass on the post-PR commit.

Term Challenge v2 will use synthetic datasets generated by swe-forge. The evaluation set is now a moving target. Every term, agents will face tasks they have never encountered before, drawn from real software projects across multiple languages. Overfitting to a static benchmark becomes impossible because the benchmark itself is continuously regenerated.

We chose to cut emissions rather than allow miners to produce useless work. We will not incentivize exploits on our network. The long-term objective is that our incentive system remains resilient to every type of vulnerability.

4 Bounty Challenge — Stricter Standards, Real Results

We have completely overhauled the incentive mechanics for Bounty Challenge.

We observed that miners were using agentic coding tools to search for bugs in the Cortex software ecosystem. This is exactly what we intended. However, the bugs reported were often similar, duplicated, or outright useless. Miners were submitting low-effort issues generated by agents without any verification, flooding the repository with noise rather than signal.

Effective immediately, the following changes apply:

Mandatory Evidence. All bug reports must include screenshots and detailed reproduction steps in the generated GitHub issues. Reports without visual evidence will be automatically rejected.

Aggressive Penalties. Invalid and duplicate issue penalties are now significantly more punitive. Each miner's valid issue count acts as a shield, but submitting invalid or duplicate bugs beyond that shield will aggressively reduce the miner's weight. The economic incentive is clear: test the software thoroughly before reporting, or lose your emissions.

Autonomous Validation Pipeline. Every bug submitted by a miner will be analyzed by an autonomous agent that validates the issue through multiple factors: reproducibility, severity, uniqueness, and relevance. Only bugs that pass this multi-factor validation will be forwarded for resolution.

12-Hour Resolution Guarantee. Every validated bug will be resolved within 12 hours. Autonomous agents will analyze the codebase, generate pull requests, and submit them for review. Each PR goes through a dual review process: a human developer manually verifies the code, and a second autonomous agent audits the PR for coherence, correctness, and potential regressions.

This pipeline transforms Bounty Challenge from a passive bug reporting system into an active, agent-powered quality assurance engine. Miners who build effective testing harnesses will be rewarded. Miners who submit noise will be penalized.

5 Cortex IDE — Alpha Release

One week ago, we announced the upcoming release of Cortex IDE. Today is the day.

We are releasing Cortex IDE in alpha. The source code is now publicly available at CortexLM/cortex-ide. This is an early release. The code is incomplete and is not yet connected to the API. This is intentional.

We are publishing the code significantly earlier than planned for a specific reason: we want to accelerate development by leveraging the miners on our network. We saw on Bounty Challenge that miners were already using agentic coding to find bugs in our software. We want to take this further. By releasing the IDE code now, we give miners a real, substantial codebase to test, audit, and improve.

Weekly Proof-of-Work Reports. Starting this week, we will publish a weekly report detailing every bug, inconsistency,

and issue that miners have found in the Cortex software products. This serves as a transparent proof of the value that decentralized miner incentives produce. The network will be able to see, in concrete terms, what the miners are contributing.

Our objective is to prove that miners have a demonstrable, measurable utility in a real software project. Not theoretical utility. Not benchmark scores. Real bugs found, real fixes shipped, real software improved.

6 Agent-Powered Development at Cortex Foundation

At Cortex Foundation, we have recently developed an internal harness that will be used across our entire organization to accelerate the production deployment of all our projects. This harness leverages autonomous agents in a structured pipeline: automated code review, automated testing, automated PR generation, with human oversight at critical checkpoints.

The architecture is proving itself. What we want on Platform Network is to maximize the performance of agents for project development. That starts today. Every tool, every pipeline, every process we build internally will feed back into the subnet's incentive system. The miners who build the best agent harnesses will directly benefit from the infrastructure we are creating.

7 On-Chain Audit and Transparency

We have deployed on the blockchain a process that will audit and report on all agent code submitted by miners. This system operates autonomously and will:

- Audit agent submissions to detect hardcoding, memorization, and other gaming strategies - Generate performance reports for the subnet, identifying what is working and what is not - Signal anomalies in the incentive system before they become exploits

The long-term objective is that Subnet 100 remains resilient to every type of vulnerability in our incentive system. We are building the tools to detect and respond to exploits before they cause damage. Transparency is not optional. Every validator, every miner, and every stakeholder will have access to the audit reports.

8 What Comes Next

Platform v2 is not an endpoint. It is the foundation for everything that follows.

In the coming weeks and months:

- Term Challenge v2 resumes with synthetic datasets from swe-forge, eliminating static benchmark gaming permanently - Bounty Challenge operates under the new strict validation and penalty regime, producing real software quality improvements - Cortex IDE development accelerates with miner-powered bug discovery and code auditing - Data Fabrication Challenge (Q2 2026) will open a new competitive frontier where miners produce coding datasets for model training - GRAIL integration will transform the produced datasets into a fine-tuned agentic coding language model

We are building the best agentic coding software on the market. We are training our own model. And we are doing it through a decentralized network where every participant is economically incentivized to contribute real value.

Platform advances. It will continue to improve.

Cortex Foundation Subnet 100, Bittensor February 20, 2026

For the full technical specification, see the Platform Network Roadmap.

References

[1] Rao, Y. et al. BitTensor: A Peer-to-Peer Intelligence Market. arXiv:2003.03917, 2020. https://arxiv.org/abs/2003.03917

[2] One Covenant Team. GRAIL: Interplanetary Intelligence. Subnet 81, 2025. https://github.com/tplr-ai/grail

[3] Jimenez, C. E. et al. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? arXiv:2310.06770, 2023. https://arxiv.org/abs/2310.06770

[4] Protocol Labs. libp2p: A Modular Network Stack. https://libp2p.io

[5] Castro, M. and Liskov, B. Practical Byzantine Fault Tolerance. OSDI, 1999. https://pmg.csail.mit.edu/papers/osdi99.pdf

[6] Haas, A. et al. Bringing the Web up to Speed with WebAssembly. PLDI, 2017. https://dl.acm.org/doi/10.1145/3140587.3062363

[7] CortexLM. swe-forge: Synthetic SWE-bench Dataset Generation. 2026. https://github.com/CortexLM/swe-forge